

ShieldAI

Financial & Client Projections — Validation & Methodology

Appendix to the Business Plan · Confidential · 2026

This appendix exists to answer the question every investor asks of an early-stage plan: **“Where do these numbers actually come from?”** Rather than assert a growth curve, we build the Year-1 client count from the bottom up — traffic, signup, conversion, and retention — and benchmark every assumption against published 2025–2026 SaaS and cybersecurity data. Where our assumptions differ from the benchmark, we explain why ShieldAI’s specific structure justifies it. The goal is not to prove the numbers are certain; it is to prove they are **defensible, conservative, and grounded in observable reality**.

1. Why These Projections Are Credible

The model does not depend on beating the market. Year-1 targets sit at or below published SaaS conversion benchmarks, and the growth engine — the managed tier — is the highest-retention, highest-margin part of the business, not the leakiest. ShieldAI’s three structural advantages (forced demand, a credit-card trial, and a managed tier with human accountability) each push our realistic outcomes *above* the SMB-SaaS median, so the plan’s use of median-or-lower assumptions is deliberately conservative.

2. Bottom-Up Build: How We Reach 170 Clients in Year 1

The headline Year-1 figure is 170 paying clients (120 Self-Serve, 30 Guided, 20 Managed). Instead of assuming that number, we derive it from a standard SaaS acquisition funnel, then show it requires only benchmark-level performance.

2.1 The Acquisition Funnel

Funnel stage	Assumption	Benchmark anchor
Qualified monthly visitors (blended organic + channel + paid)	~2,500 / mo by end of Y1 ramp	Modest for a funded launch with insurer/MSP referral channels
Visitor → trial signup	~7%	Cybersecurity visitor-to-trial ~7.4% (SaaSHero, 2026)
Trial → paid conversion	~22%	Cybersecurity trial-to-paid 21.9% (First Page Sage); CC-required trials 30%+ (ChartMogul, 2026)
Resulting new paid / month (steady state)	~35–40	= 2,500 × 7% × 22%
Net Year-1 paying clients after ramp & churn	~170	Ramp from a standing start; see §3 for churn

The funnel is intentionally undersized: 2,500 qualified visitors/month is achievable through content, insurer/MSP referrals, and modest paid spend within the go-to-market budget. We do not assume virality or outsized paid efficiency.

2.2 Why Our Conversion Assumptions Are Conservative, Not Optimistic

- **Cybersecurity converts above the SaaS median.** Trial-to-paid in cybersecurity runs ~21.9%, versus an ~18.5% B2B-SaaS median — problem urgency lifts conversion. We model 22%, essentially at the category benchmark, not above it.

- **A credit-card (opt-out) trial more than doubles conversion.** ChartMogul's 2026 study of 200 products found CC-required trials convert at 30%+ — over 5× opt-in trials. If ShieldAI requires a card at trial (recommended), 22% is well below what that structure typically yields, giving material upside we do not bank on.
- **Forced demand shortens the “why now.”** Insurance, customer questionnaires, and regulation compel action on a deadline. Trust signals (compliance badges, security certifications) lift FinTech/security conversions up to 40% (M Accelerator, 2025) — a lever directly available to us.

3. Retention: The Assumption That Usually Breaks SMB Models — and Why Ours Holds

The single biggest risk to any SMB-SaaS projection is churn. Self-serve SMB products routinely churn 3–7% per month (roughly 31–58% per year). A plan that ignores this is not credible. Ours confronts it directly and is structured specifically to beat it.

Segment	Published monthly churn	What we assume	Why
Self-Serve (\$129)	3–7% (SMB self-serve)	~5% (mid-range)	Priced/positioned as SMB self-serve; we do not assume best-in-class here
Guided (\$599)	2–4% (mid-market-like)	~3%	Periodic human review raises switching cost and perceived value
Managed (\$1,950)	0.5–2% (enterprise-like)	~1.5%	Human-run program, compliance cycles, deep integration — enterprise-grade stickiness

Sources: Recurly 2025 Churn Report; Optifai 2026 (N=939); SaaS Capital 2025. SMB self-serve 3–7%/mo; enterprise 0.5–2%/mo. Higher-ARPU tiers retain dramatically better — the exact dynamic our tier ladder is built to exploit.

The strategic point: our growth and revenue concentrate in the managed tier, which is the *stickiest* part of SMB software (enterprise-like ~1.5%/mo), not the leakiest. Self-serve is the low-cost wedge that feeds it. This is why blended economics improve as the mix shifts to managed — the opposite of a typical churn-fragile SMB model.

3.1 The Expansion Engine (Net Revenue Retention)

Because clients enter at \$129 and graduate to \$599 and \$1,950, expansion revenue partially offsets logo churn. Best-in-class SaaS reaches 110%+ net revenue retention through exactly this kind of tier-graduation loop (SaaS Capital, 2025). We do not model negative churn, but the land-and-expand ladder is a structural tailwind that our conservative per-tier churn assumptions **ignore** — additional cushion, not additional risk.

4. Unit Economics: Stress-Testing the Per-EngineER Claim

The investment thesis rests on one number: a single engineer, amplified by the console and AI co-pilot, profitably manages ~25 managed clients. We stress-test it below and show the model still works even if that number is materially worse.

Scenario	Clients / engineer	Rev / engineer / yr	Gross contribution*	Still viable?
Plan case	25	\$585,000	~\$440,000 (~75%)	Yes
Conservative	18	\$421,000	~\$276,000 (~66%)	Yes
Pessimistic	12	\$281,000	~\$136,000 (~48%)	Yes — still

Scenario	Clients / engineer	Rev / engineer / yr	Gross contribution*	Still viable?
				positive

*Gross contribution after ~\$140K fully-loaded engineer cost and ~\$15/client/month AI + infrastructure. Even at 12 clients/engineer — less than half the plan assumption — the managed tier stays gross-margin positive. The model is robust to being wrong about its central efficiency claim.

5. Triangulation: Three Independent Paths to the Same Conclusion

A projection is stronger when independent methods agree. Our Year-3 target (~\$10M ARR, ~1,600 clients) is corroborated three ways:

- **Bottom-up funnel (\$2):** benchmark conversion on a modest, growing traffic base compounds to this range over three years with mix shift toward managed.
- **Top-down market share:** ~1,600 clients is a rounding error — roughly 0.03% of the ~5.8M U.S. employer businesses, and well under 1% of even the conservative 600K–1.2M regulated/insured SAM. The plan requires trivial penetration.
- **Capacity check:** ~11 engineers at 25 managed clients each supports ~275 managed clients — consistent with the Year-3 managed count. Headcount, revenue, and funnel all reconcile.

When a bottom-up funnel, a top-down market-share sanity check, and a delivery-capacity model all land in the same place, the projection is not a guess — it is a triangulated estimate. The risk in the plan is execution speed, not arithmetic.

6. What Would Make Us Wrong (and How We De-Risk It)

If we are wrong about...	The evidence would show...	Mitigation already in the plan
Self-serve churn	Self-serve >7%/mo	Managed tier (not self-serve) drives revenue; self-serve is a low-CAC feeder
Conversion rate	Trial→paid <15%	CC-required trial + insurer/MSP referral channels raise both quality and rate
Engineer efficiency	<12 clients/engineer	Model stays margin-positive at 12; console/AI investment is the direct lever
CAC	Paid CAC too high for \$129 tier	Channel/referral motion (insurers, MSPs) is low-CAC by design; self-serve is wedge, not profit center

Every primary risk maps to a mitigation that is structural — built into the tier design and go-to-market — rather than a hope that the market behaves better than benchmark.

7. Summary

These projections are conservative by construction. Conversion is modeled at the cybersecurity benchmark despite a credit-card trial that typically beats it. Churn is modeled at mid-range for self-serve and enterprise-like only where the human-managed model earns it. The revenue engine is the stickiest tier, not the leakiest. The central efficiency claim survives being cut in half. And three independent methods converge on the same Year-3 range, which itself requires under 1% of the serviceable market.

The plan does not need optimism to work. It needs execution against assumptions that published 2025–2026 data already supports — which is exactly the profile an investor should want at this stage, especially paired with a product that is already live in production.

Benchmark sources: First Page Sage (2025), ChartMogul SaaS Conversion Report (2026), Recurly Churn Report (2025), Optifai (2026, N=939), SaaS Capital (2025), SaaSHero (2026), M Accelerator (2025), SBA Office of Advocacy (2025). Figures are planning estimates for illustration, not guarantees.