

ShieldAI

The Virtual CISO Platform

● LIVE IN PRODUCTION · Executive Summary · Confidential · 2026

“A full-time CISO costs \$200,000–\$400,000 a year. Nearly half of small businesses have no security budget at all — yet their insurers, customers, and regulators now require a real security program. ShieldAI closes that gap — and it’s already live.”

ShieldAI is a live, in-production virtual Chief Information Security Officer (vCISO) platform for small and mid-sized businesses. It is not a concept, a prototype, or a set of mockups — it is deployed, and anyone can create an account and use it today. A visitor can register, run a real security assessment, receive a deterministic NIST-based risk score, and generate written policies and a prioritized remediation roadmap in minutes.

The platform delivers a complete, continuously-managed security program — risk scoring, written policies, a prioritized roadmap, compliance mapping, awareness training, and executive reporting — at a price a small business can actually afford. Behind it runs a production stack: a multi-provider AI layer (Claude, Gemini, and GPT), a tiered AI advisor scoped to each customer’s plan, and live threat intelligence drawn from real CVE and breach-exposure data.

Strategy: Wedge to Platform

Our strategy is a deliberate wedge-to-platform progression. We begin with the affordable, largely self-serve product that is live today and delivers immediate value. As we grow, we layer human expertise on top: cybersecurity engineers who use ShieldAI to run full, managed CISO programs for a portfolio of clients — with the software and its AI co-pilot doing the heavy lifting so a small team can protect a large book of businesses. This is how we move from a low-cost tool to a high-value managed-security platform without a linear increase in headcount.

Where the Moat Actually Is

ShieldAI is powered by frontier AI for analysis and content generation, but our defensibility does not rest on the AI itself. It rests on three things: a proprietary, deterministic security-scoring engine built on the NIST Cybersecurity Framework that produces explainable, auditor- and insurer-trustworthy results; a human-in-the-loop model in which qualified engineers review and approve work before it reaches the client; and the proprietary assessment data and client relationships that compound over time. The AI is a commodity input — the scoring system, the human review layer, and the accumulating data are the moat.

Key Highlights

Status	LIVE IN PRODUCTION — deployed, and usable today at shieldai-production-627e.up.railway.app
Target market	U.S. SMBs with employees, prioritizing regulated and insured segments
Core category	Virtual CISO (vCISO) — ~\$2B+ today, growing 12–15%+ annually
Demand driver	Cyber-insurance, customer security requirements, and regulation forcing SMBs to act

Model	Wedge: low-cost self-serve → Expand: engineer-run managed vCISO programs
Differentiation	Deterministic NIST scoring engine + human-in-the-loop + multi-model AI + analyst platform
Built & working	Assessment, deterministic scoring, AI program & policy generation, tiered Mastermind advisor, live CVE/breach threat intel, monitoring agents, accounts, billing, admin
Next milestone	Scale the managed-services layer: onboard first paying clients and hire the first security engineers

Most companies at our stage are pitching a plan. We are pitching a product that already exists and works. The self-serve platform is built, deployed, and functioning. This raise funds the build-out of the managed-services layer that turns a working product into a scalable managed-security business — and lets a handful of our engineers protect thousands of small businesses.